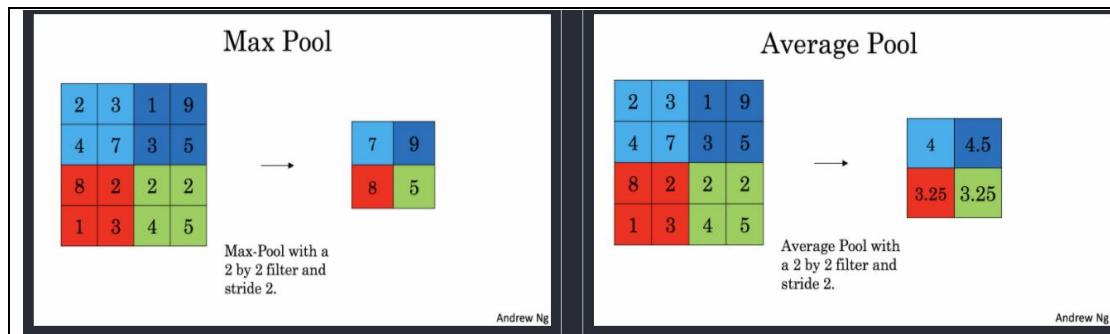


计算机科学与技术学院神经网络与深度学习课程实验报告

实验题目：Convolutional Neural Networks and ResNets		学号：201900130151
日期：2021. 11. 2	班级： 人工智能	姓名： 莫甫龙
Email：m1533979510@163. com		
实验目的： 了解卷积神经网络的结构网络 了解 Resnet 网络		
实验软件和硬件环境： Vs code Win11		
实验原理和方法： <pre>a = np.pad(a, ((0,0), (1,1), (0,0), (3,3), (0,0)), 'constant', constant_values = (...))</pre> 使用 np. pad 函数来使用 0 扩充边界。 卷积神经网络的前向传播： <pre>when the input shape is:</pre> $n_H = \lfloor \frac{n_{H_{prev}} - f + 2 \times pad}{stride} \rfloor + 1$ $n_W = \lfloor \frac{n_{W_{prev}} - f + 2 \times pad}{stride} \rfloor + 1$ $n_C = \text{number of filters used in the convolution}$ 池化层的前向传播： <pre>for the pooling to the input shape is:</pre> $n_H = \lfloor \frac{n_{H_{prev}} - f}{stride} \rfloor + 1$ $n_W = \lfloor \frac{n_{W_{prev}} - f}{stride} \rfloor + 1$ $n_C = n_{C_{prev}}$		



最大值池化层：在输入矩阵中滑动一个大小为 $f \times f$ 的窗口，选取窗口里的值中的最大值，然后作为输出。

均值池化层：在输入矩阵中滑动一个大小为 $f \times f$ 的窗口，计算窗口中所有值的平均值，然后作为输出。


卷积层的反向传播：

$$dA_{+} = \sum_{h=0}^{n_H} \sum_{w=0}^{n_W} W_c \times dZ_{hw}$$

$$dW_{c+} = \sum_{h=0}^{n_H} \sum_{w=0}^{n_W} a_{slice} \times dZ_{hw}$$

$$db = \sum_h \sum_w dZ_{hw}$$

Resnet



实验步骤：（不要求罗列完整源代码）

Convolutional Neural Networks: Step by Step:

Zero_pad:

```
### START CODE HERE ### (~ 1 line)
X_pad = np.pad(X, ((0,0),(pad,pad),(pad,pad),(0,0)), 'constant')
### END CODE HERE ###
```

conv_single_step:

```
# Element-wise product between a_slice and w. Add bias
s = w*a_slice_prev
# Sum over all entries of the volume s
Z = np.sum(s)+b
### END CODE HERE ###
```

conv_forward:

```

# Retrieve dimensions from A_prev's shape (~1 line)
(m, n_H_prev, n_W_prev, n_C_prev) = A_prev.shape

# Retrieve dimensions from W's shape (~1 line)
(f, f, n_C_prev, n_C) = W.shape

# Retrieve information from "hparameters" (~2 lines)
stride = hparameters['stride']
pad = hparameters['pad']

# Compute the dimensions of the CONV output volume using the formula given above. Hint: use int() to floor. (~2 lines)
n_H = int((n_H_prev-f+2*pad)/stride+1)
n_W = int((n_W_prev-f+2*pad)/stride+1)

# Initialize the output volume Z with zeros. (~1 line)
Z = np.zeros((m,n_H,n_W,n_C))

# Create A_prev_pad by padding A_prev
A_prev_pad = zero_pad(A_prev, pad)

```

```

for i in range(m):
    a_prev_pad = A_prev_pad[i]
    for h in range(n_H):
        for w in range(n_W):
            for c in range(n_C):
                # Loop over the batch of training examples
                # Select ith training example's padded activation
                # Loop over vertical axis of the output volume
                # Loop over horizontal axis of the output volume
                # Loop over channels (= #filters) of the output volume

                # Find the corners of the current "slice" (~4 lines)
                vert_start = h * stride
                vert_end = vert_start + f
                horiz_start = w * stride
                horiz_end = horiz_start + f

                # Use the corners to define the (3D) slice of a_prev_pad (See Hint above the cell). (~1 line)
                a_slice_prev = a_prev_pad[vert_start:vert_end,horiz_start:horiz_end,:]

                # Convolve the (3D) slice with the correct filter W and bias b, to get back one output neuron. (~1 line)
                Z[i, h, w, c] = conv_single_step(a_slice_prev, W[:, :, :, c], b[:, :, :, c])

### END CODE HERE ###

# Making sure your output shape is correct
assert(Z.shape == (m, n_H, n_W, n_C))

# Save information in "cache" for the backprop
cache = (A_prev, W, b, hparameters)

return Z, cache

```

pool_forward:

```

# Retrieve dimensions from the input shape
(m, n_H_prev, n_W_prev, n_C_prev) = A_prev.shape

# Retrieve hyperparameters from "hparameters"
f = hparameters["f"]
stride = hparameters["stride"]

# Define the dimensions of the output
n_H = int(1 + (n_H_prev - f) / stride)
n_W = int(1 + (n_W_prev - f) / stride)
n_C = n_C_prev

# Initialize output matrix A
A = np.zeros((m, n_H, n_W, n_C))

```

```

### START CODE HERE ###
for i in range(m):
    for h in range(n_H):
        for w in range(n_W):
            for c in range(n_C):
                # Loop over the training examples
                # Loop on the vertical axis of the output volume
                # Loop on the horizontal axis of the output volume
                # Loop over the channels of the output volume

                # Find the corners of the current "slice" (~4 lines)
                vert_start = h * stride
                vert_end = vert_start + f
                horiz_start = w * stride
                horiz_end = horiz_start + f

                # Use the corners to define the current slice on the ith training example of A_prev, channel c. (~1 line)
                a_prev_slice = A_prev[i, vert_start:vert_end, horiz_start:horiz_end, c]

                # Compute the pooling operation on the slice. Use an if statement to differentiate the modes. Use np.max/np.mean.
                if mode == "max":
                    A[i, h, w, c] = np.max(a_prev_slice)
                elif mode == "average":
                    A[i, h, w, c] = np.mean(a_prev_slice)

### END CODE HERE ###

# Store the input and hparameters in "cache" for pool_backward()
cache = (A_prev, hparameters)

# Making sure your output shape is correct
assert(A.shape == (m, n_H, n_W, n_C))

```

conv_backward:

```

# Retrieve information from "cache"
(A_prev, W, b, hparameters) = cache

# Retrieve dimensions from A_prev's shape
(m, n_H_prev, n_W_prev, n_C_prev) = A_prev.shape

# Retrieve dimensions from W's shape
(f, f, n_C_prev, n_C) = W.shape

# Retrieve information from "hparameters"
stride = hparameters['stride']
pad = hparameters['pad']

# Retrieve dimensions from dZ's shape
(m, n_H, n_W, n_C) = dZ.shape

# Initialize dA_prev, dW, db with the correct shapes
dA_prev = np.zeros_like(A_prev)
dW = np.zeros_like(W)
db = np.zeros_like(b)

# Pad A_prev and dA_prev
A_prev_pad = zero_pad(A_prev, pad)
dA_prev_pad = zero_pad(dA_prev, pad)

```

```

for i in range(m):                # loop over the training examples

    # select ith training example from A_prev_pad and dA_prev_pad
    a_prev_pad = A_prev_pad[i]
    da_prev_pad = dA_prev_pad[i]

    for h in range(n_H):          # loop over vertical axis of the output volume
        for w in range(n_W):      # loop over horizontal axis of the output volume
            for c in range(n_C):  # loop over the channels of the output volume

                # Find the corners of the current "slice"
                vert_start = h * stride
                vert_end = vert_start + f
                horiz_start = w * stride
                horiz_end = horiz_start + f

                # Use the corners to define the slice from a_prev_pad
                a_slice = a_prev_pad[vert_start:vert_end, horiz_start:horiz_end, :]

                # Update gradients for the window and the filter's parameters using the code formulas given
                da_prev_pad[vert_start:vert_end, horiz_start:horiz_end, :] += W[:, :, :, c] * dz[i, h, w, c]
                dW[:, :, :, c] += a_slice * dz[i, h, w, c]
                db[:, :, :, c] += dz[i, h, w, c]

            # Set the ith training example's dA_prev to the unpadded da_prev_pad (Hint: use X[pad:-pad, pad:-pad, :])
            dA_prev[i, :, :, :] = da_prev_pad[pad:-pad, pad:-pad, :]
        ### END CODE HERE ###

    # Making sure your output shape is correct
    assert(dA_prev.shape == (m, n_H_prev, n_W_prev, n_C_prev))

    return dA_prev, dW, db

```

create_mask_from_window:

```

    """ START CODE HERE """ (~1 line)
    mask = (np.max(x)==x)
    """ END CODE HERE """

    return mask

```

distribute_value:

```

    """ START CODE HERE """
    # Retrieve dimensions from shape (~1 line)
    (n_H, n_W) = shape

    # Compute the value to distribute on the matrix (~1 line)
    average = dz/(n_H*n_W)

    # Create a matrix where every entry is the "average" value (~1 line)
    a = np.ones(shape) * average
    """ END CODE HERE """

    return a

```

pool_backward:

```

    # Retrieve information from cache (~1 line)
    (A_prev, hparameters) = cache

    # Retrieve hyperparameters from "hparameters" (~2 lines)
    stride = hparameters['stride']
    f = hparameters['f']

    # Retrieve dimensions from A_prev's shape and dA's shape (~2 lines)
    m, n_H_prev, n_W_prev, n_C_prev = A_prev.shape
    m, n_H, n_W, n_C = dA.shape

    # Initialize dA_prev with zeros (~1 line)
    dA_prev = np.zeros_like(A_prev)

```

```

for i in range(m):                                # Loop over the training examples

    # select training example from A_prev (~1 line)
    a_prev = A_prev[i]

    for h in range(n_H):                          # Loop on the vertical axis
        for w in range(n_W):                      # Loop on the horizontal axis
            for c in range(n_C):                  # Loop over the channels (depth)

                # Find the corners of the current "slice" (~4 lines)
                vert_start = h * stride
                vert_end = vert_start + f
                horiz_start = w * stride
                horiz_end = horiz_start + f

                # Compute the backward propagation in both modes.
                if mode == "max":

                    # Use the corners and "c" to define the current slice from a_prev (~1 line)
                    a_prev_slice = a_prev[vert_start:vert_end, horiz_start:horiz_end, c]
                    # Create the mask from a_prev_slice (~1 line)
                    mask = create_mask_from_window(a_prev_slice)
                    # Set dA_prev to be dA_prev + (the mask multiplied by the correct entry of dA) (~1 line)
                    dA_prev[i, vert_start: vert_end, horiz_start: horiz_end, c] += np.multiply(mask, dA[i,h,w,c])

                elif mode == "average":

                    # Get the value a from dA (~1 line)
                    da = dA[i,h,w,c]
                    # Define the shape of the filter as fxf (~1 line)
                    shape = (f,f)
                    # Distribute it to get the correct slice of dA_prev. i.e. Add the distributed value of da. (~1 line)
                    dA_prev[i, vert_start: vert_end, horiz_start: horiz_end, c] += distribute_value(da, shape)

```

Convolutional Neural Networks: Application:

create_placeholders:

```

### START CODE HERE ### (~2 lines)
X = tf.placeholder(tf.float32, shape=(None, n_H0, n_W0, n_C0))
Y = tf.placeholder(tf.float32, shape=(None, n_y))
### END CODE HERE ###

return X, Y

```

initialize_parameters:

```

tf.set_random_seed(1)                                # so that your "random" numbers match ours

### START CODE HERE ### (approx. 2 lines of code)
W1 = tf.get_variable("W1", [4, 4, 3, 8], initializer=tf.contrib.layers.xavier_initializer(seed=0))
W2 = tf.get_variable("W2", [2, 2, 8, 16], initializer=tf.contrib.layers.xavier_initializer(seed=0))
### END CODE HERE ###

parameters = {"W1": W1,
              "W2": W2}

```

forward_propagation:


```

# Retrieve the parameters from the dictionary "parameters"
W1 = parameters['W1']
W2 = parameters['W2']

### START CODE HERE ###
# CONV2D: stride of 1, padding 'SAME'
Z1 = tf.nn.conv2d(X, W1, strides=(1, 1, 1, 1), padding='SAME')
# RELU
A1 = tf.nn.relu(Z1)
# MAXPOOL: window 8x8, stride 8, padding 'SAME'
P1 = tf.nn.max_pool(A1, ksize=(1, 8, 8, 1), strides=(1, 8, 8, 1), padding='SAME')
# CONV2D: filters W2, stride 1, padding 'SAME'
Z2 = tf.nn.conv2d(P1, W2, strides=(1, 1, 1, 1), padding='SAME')
# RELU
A2 = tf.nn.relu(Z2)
# MAXPOOL: window 4x4, stride 4, padding 'SAME'
P2 = tf.nn.max_pool(A2, ksize=(1, 4, 4, 1), strides=(1, 4, 4, 1), padding='SAME')
# FLATTEN
P2 = tf.contrib.layers.flatten(P2)
# FULLY-CONNECTED without non-linear activation function (not not call softmax).
# 6 neurons in output layer. Hint: one of the arguments should be "activation_fn=None"
Z3 = tf.contrib.layers.fully_connected(P2, num_outputs = 6, activation_fn=None)
### END CODE HERE ###

```

compute_cost:

```

### START CODE HERE ### (1 line of code)
cost = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits = Z3, labels = Y))
### END CODE HERE ###

```

model:

```

X, Y = create_placeholders(n_H0, n_W0, n_C0, n_y)
### END CODE HERE ###

# Initialize parameters
### START CODE HERE ### (1 line)
parameters = initialize_parameters()
### END CODE HERE ###

# Forward propagation: Build the forward propagation in the tensorflow graph
### START CODE HERE ### (1 line)
Z3 = forward_propagation(X, parameters)
### END CODE HERE ###

# Cost function: Add cost function to tensorflow graph
### START CODE HERE ### (1 line)
cost = compute_cost(Z3, Y)
### END CODE HERE ###

# Backpropagation: Define the tensorflow optimizer. Use an AdamOptimizer that minimizes
### START CODE HERE ### (1 line)
optimizer = tf.train.AdamOptimizer(learning_rate = learning_rate).minimize(cost)
### END CODE HERE ###

# Initialize all the variables globally
init = tf.global_variables_initializer()

```

```

# Start the session to compute the tensorflow graph
with tf.Session() as sess:

    # Run the initialization
    sess.run(init)

    # Do the training loop
    for epoch in range(num_epochs):

        minibatch_cost = 0.
        num_minibatches = int(m / minibatch_size) # number of minibatches of size minibatch_size in the train set
        seed = seed + 1
        minibatches = random_mini_batches(X_train, Y_train, minibatch_size, seed)

        for minibatch in minibatches:

            # Select a minibatch
            (minibatch_X, minibatch_Y) = minibatch
            # IMPORTANT: The line that runs the graph on a minibatch.
            # Run the session to execute the optimizer and the cost, the feeddict should contain a minibatch for (
            ### START CODE HERE ### (1 line)
            _, temp_cost = sess.run([optimizer, cost], feed_dict={X: minibatch_X, Y: minibatch_Y})
            ### END CODE HERE ###

            minibatch_cost += temp_cost / num_minibatches

        # Print the cost every epoch
        if print_cost == True and epoch % 5 == 0:
            print ("Cost after epoch %i: %f" % (epoch, minibatch_cost))
        if print_cost == True and epoch % 1 == 0:
            costs.append(minibatch_cost)

```

Resnet:

identity_block:

```

# Second component of main path (~3 lines)
X = Conv2D(filters = F2, kernel_size = (f, f), strides = (1,1), padding = 'same', name = conv_name_base + '2b', kernel_initializer = glorot_uniform)
X = BatchNormalization(axis = 3, name = bn_name_base + '2b')(X)
X = Activation('relu')(X)

# Third component of main path (~2 lines)
X = Conv2D(filters = F3, kernel_size = (1, 1), strides = (1,1), padding = 'valid', name = conv_name_base + '2c', kernel_initializer = glorot_uniform)
X = BatchNormalization(axis = 3, name = bn_name_base + '2c')(X)

# Final step: Add shortcut value to main path, and pass it through a RELU activation (~2 lines)
X = Add()(X, X_shortcut))
X = Activation('relu')(X)

### END CODE HERE ###

```

convolutional_block:

```

# Second component of main path (~3 lines)
X = Conv2D(filters = F2, kernel_size = (f, f), strides = (1,1), padding = 'same', name = conv_name_base + '2b', kernel_initializer = glorot_uniform)
X = BatchNormalization(axis = 3, name = bn_name_base + '2b')(X)
X = Activation('relu')(X)

# Third component of main path (~2 lines)
X = Conv2D(filters = F3, kernel_size = (1, 1), strides = (1,1), padding = 'valid', name = conv_name_base + '2c', kernel_initializer = glorot_uniform)
X = BatchNormalization(axis = 3, name = bn_name_base + '2c')(X)

##### SHORTCUT PATH ##### (~2 lines)
X_shortcut = Conv2D(filters = F3, kernel_size = (1, 1), strides = (s,s), padding = 'valid', name = conv_name_base + '1', kernel_initializer = glorot_uniform)
X_shortcut = BatchNormalization(axis = 3, name = bn_name_base + '1')(X_shortcut)

# Final step: Add shortcut value to main path, and pass it through a RELU activation (~2 lines)
X = Add()(X, X_shortcut))
X = Activation('relu')(X)

### END CODE HERE ###

```

ResNet50:

```
### START CODE HERE ###

# Stage 3 (~4 lines)
X = convolutional_block(X, f = 3, filters = [128, 128, 512], stage = 3, block='a', s = 2)
X = identity_block(X, 3, [128, 128, 512], stage=3, block='b')
X = identity_block(X, 3, [128, 128, 512], stage=3, block='c')
X = identity_block(X, 3, [128, 128, 512], stage=3, block='d')

# Stage 4 (~6 lines)
X = convolutional_block(X, f = 3, filters = [256, 256, 1024], stage = 4, block='a', s = 2)
X = identity_block(X, 3, [256, 256, 1024], stage=4, block='b')
X = identity_block(X, 3, [256, 256, 1024], stage=4, block='c')
X = identity_block(X, 3, [256, 256, 1024], stage=4, block='d')
X = identity_block(X, 3, [256, 256, 1024], stage=4, block='e')
X = identity_block(X, 3, [256, 256, 1024], stage=4, block='f')

# Stage 5 (~3 lines)
X = convolutional_block(X, f = 3, filters = [512, 512, 2048], stage = 5, block='a', s = 2)
X = identity_block(X, 3, [512, 512, 2048], stage=5, block='b')
X = identity_block(X, 3, [512, 512, 2048], stage=5, block='c')

# AVGPOOL (~1 line). Use "X = AveragePooling2D(...)(X)"
X = AveragePooling2D(pool_size=(2, 2))(X)

### END CODE HERE ###
```

结论分析与体会:

在 TensorFlow 里面有一些可以直接拿来用的函数可以直接实现要实现的功能, 这个在实际应用中会方便很多。

Resnet 的出现, 让深度网络模型成为了可能, 让深度不再局限于 20 层。

就实验过程中遇到和出现的问题, 你是如何解决和处理的, 自拟 1—3 道问答题:

1. 在装 tensorflow 环境的时候, 因为没有考虑到 python 和 tensorflow 之间的版本依赖问题, 导致配的环境有问题, 有些 tensorflow1 的函数在 tensorflow2 中不能运行 (因为已经被删了)。